
Practical Two-Step Look-Ahead Bayesian Optimization

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Expected improvement and other acquisition functions widely used in Bayesian
2 optimization use a “one-step” assumption: they value objective function evalu-
3 ations assuming no future evaluations will be performed. Because we usually
4 evaluate over multiple steps, this assumption may leave substantial room for im-
5 provement. Existing theory gives acquisition functions looking multiple steps in the
6 future but calculating them requires solving a high-dimensional continuous-state
7 continuous-action Markov decision process (MDP). Fast exact solutions of this
8 MDP remain out of reach of today’s methods. As a result, previous two- and
9 multi-step lookahead Bayesian optimization algorithms are either too expensive to
10 implement in most practical settings or resort to heuristics that may fail to fully
11 realize the promise of two-step lookahead. This paper proposes a computationally
12 efficient algorithm that provides an accurate solution to the two-step lookahead
13 Bayesian optimization problem in seconds to at most several minutes of computa-
14 tion per batch of evaluations, and provides increased query efficiency compared
15 with previous two- and multi-step lookahead methods. This unlocks the value of
16 two-step lookahead in practice. We demonstrate the value of our algorithm with
17 extensive experiments on synthetic test functions and real-world problems.

18 1 Introduction

19 We consider minimization of a continuous black-box function f over a hyperrectangle $\mathcal{A} \subseteq \mathbb{R}^d$.
20 We suppose evaluations $f(x)$ are time-consuming to obtain, do not provide first- or second-order
21 derivative information and noise-free. Such problems arise when tuning hyperparameters of complex
22 machine learning models [Snoek et al., 2012] and optimizing engineering systems using physics-based
23 simulators [Forrester et al., 2008].

24 We consider this problem within a Bayesian optimization (BayesOpt) framework [Brochu et al.,
25 2010]. BayesOpt methods contain two components: (1) a statistical model over f , typically a
26 Gaussian process [Rasmussen and Williams, 2006]; and (2) an acquisition function computed from
27 the statistical model that quantifies the value of evaluating f . After a first stage of evaluations of
28 f , often at points chosen uniformly at random from \mathcal{A} , we behave iteratively: we fit the statistical
29 model to all available data; then optimize the resulting acquisition function (which can be evaluated
30 quickly and often provides derivative information) to find the best point(s) to evaluate f ; perform
31 these evaluations; and repeat until our evaluation budget is exhausted.

32 The most widely-used acquisition functions use a one-step lookahead approach. They consider the
33 *direct* effect of the evaluation on an immediate measure of solution quality, and do not consider
34 evaluations that will be performed later. This includes expected improvement (EI) [Jones et al., 1998],
35 probability of improvement (PI) Kushner [1964], entropy search (ES) [Hernández-Lobato et al., 2014],
36 Wang and Jegelka, 2017], and the knowledge gradient (KG) [Wu and Frazier, 2016]. By myopically

37 maximizing the immediate improvement in solution quality, they may sacrifice even greater gains in
38 solution quality obtainable through coordinated action across multiple evaluations.

39 Researchers have sought to address this shortcoming through non-myopic acquisition functions. The
40 decision of where to sample next in BayesOpt can be formulated as a partially observable Markov
41 decision process (POMDP) [Ginsbourger and Riche, 2010]. The solution to this POMDP is given by
42 the Bellman recursion [Lam et al., 2016] and provides a non-myopic acquisition function that provides
43 the best possible average-case performance under the prior. However, the “curse of dimensionality”
44 Powell [2007] prevents solving this POMDP for even small-scale problems.

45 The past literature [Lam et al., 2016, Osborne et al., 2009, Ginsbourger and Riche, 2010, González
46 et al., 2016] instead approximates the solution to this POMDP to create non-myopic acquisition
47 functions. Two-step lookahead is particularly attractive [Osborne et al., 2009, Ginsbourger and Riche,
48 2010, González et al., 2016] because it is substantially easier to compute than looking ahead more than
49 two steps, but still promises a performance improvement over the one-step acquisition functions used
50 in practice. Indeed, Ginsbourger and Riche [2010] argue that using two-step lookahead encourages
51 a particularly beneficial form of exploration: evaluating a high uncertainty region benefits future
52 evaluations; if the evaluation reveals the region was better than expected, then future evaluations
53 evaluate nearby to find improvements in solution quality. This benefit occurs even if the first
54 evaluation does not generate a direct improvement in solution quality. In numerical experiments,
55 Osborne et al. [2009], Ginsbourger and Riche [2010] show that two-step lookahead improves over
56 one-step lookahead in a range of practical problems.

57 At the same time, optimizing two-step acquisition functions is computationally challenging. Unlike
58 common one-step acquisition functions like expected improvement, they cannot be computed in
59 closed form and instead require a time-consuming simulation with nested optimization. Simulation
60 creates noise and prevents straightforward differentiation, which hampers optimizing these two-
61 step acquisition functions precisely. Existing approaches [Osborne et al., 2009, Ginsbourger and
62 Riche, 2010, González et al., 2016] use derivative-free optimizers, which can require a large number
63 of iterations to optimize precisely, particularly as the dimension d of the feasible space grows.
64 (Numerical experiments in Osborne et al. [2009], Ginsbourger and Riche [2010] are restricted to
65 problems with $d \leq 3$.) As a result, existing two- and multi-step methods require a prohibitive amount
66 of computation (e.g., Lam [2018] reports that the method in Lam et al. [2016] requires between 10
67 minutes and 1 hour per evaluation even on low-dimensional problems). If sufficient computation is
68 not performed, then errors in the acquisition-function optimization overwhelm the benefits provided
69 by two-step lookahead and query efficiency degrades compared to a one-step acquisition function
70 supporting precise optimization. Similar challenges arise for the multi-step method proposed in
71 Lam et al. [2016]. This computational challenge has largely prevented the widespread adoption of
72 non-myopic acquisition functions in practice.

73 **Contributions.** This article makes two key innovations unlocking the power of the two-step looka-
74 head in practice. First, we provide an unbiased estimator based on the envelope theorem for the
75 gradient of the two-step lookahead acquisition function, which can be used within a stochastic opti-
76 mizer. This supports optimizing the acquisition function accurately with computation time ranging
77 between a few seconds and several minutes, and scales better in the batch size and dimension of the
78 black-box function compared with the common practice of using a derivative-free optimizer. Second,
79 we embrace the variance reduction in Monte-Carlo simulations to further reduce the computational
80 cost of estimating both the two-step lookahead acquisition function and its gradient. We prove that
81 our gradient estimator is unbiased consistent and the stochastic gradient ascent used for optimizing
82 the acquisition function will converge to a local stationary point. We will make our implementation
83 and the code to replicate the experiments public upon publication.

84 Our approach leverages computational techniques developed in the literature. The first is the use of
85 infinitesimal perturbation analysis [Heidelberger et al., 1988] and the envelope theorem [Milgrom
86 and Segal, 2002], previously used in Bayesian optimization to optimize the knowledge gradient
87 acquisition function (which is myopic, as noted above) by Wu et al. [2017]. The second is a pair
88 of variance reduction methods: Gauss-Hermite quadrature Liu and Pierce [1994] and importance
89 sampling [Asmussen and Glynn, 2007]. Our paper is the first to demonstrate the power of these
90 techniques for non-myopic Bayesian optimization.

2 The Two-Step Optimal (2-OPT) Acquisition Function

This section defines the two-step lookahead acquisition function. This acquisition function is optimal when there are two stages of measurements remaining, and so we call it 2-OPT. Before defining 2-OPT, we first provide notation and brief background from Gaussian process regression in Sect. 2.1. We then define 2-OPT in Sect. 2.2 and show how to estimate it with Monte Carlo in Sect. 2.3. While 2-OPT has been defined implicitly in past work, we include a complete description to provide a framework and notation supporting our novel efficient method for maximizing 2-OPT in Sect. 3.

2.1 Gaussian process model for the objective f

We place a Gaussian process (GP) prior on the objective f . Although standard, here we briefly describe inference under a GP to provide notation used later. Our GP prior is characterized by a mean function $\mu(\cdot)$ and a kernel function $K(\cdot, \cdot)$. The posterior distribution on f after observing f at data points $D = (x^{(1)}, \dots, x^{(m)})$ is a GP with mean function and kernel defined respectively by

$$\begin{aligned} \mu(x) + K(x, D)K(D, D)^{-1}(f(D) - \mu(D)), \\ K(x, x') - K(x, D)K(D, D)^{-1}K(D, x'). \end{aligned} \quad (1)$$

In (1), $f(D) = (f(x^{(1)}), \dots, f(x^{(m)}))$, and similarly for $\mu(D)$. Expressions $K(x, D)$, $K(D, x)$, and $K(D, D)$ similarly evaluate to a column vector, row vector, and square matrix respectively.

2.2 Two-step lookahead acquisition function

Here we define the 2-OPT acquisition function from a theoretical (but not yet computational) perspective. This formulation follows previous work on two-step and multi-step acquisition functions [Lam et al., 2016, Osborne et al., 2009, Ginsbourger and Riche, 2010, González et al., 2016]. 2-OPT gives optimal average-case behavior when we have two stages of evaluations remaining, and the second stage of evaluation may be chosen based on the results from the first.

To support batch evaluations while maintaining computational tractability, our first stage of evaluations uses a batch of $q \geq 1$ simultaneous evaluations, while the second stage uses a single evaluation.

Throughout, we assume that we have already observed a collection of data points D , so that the current posterior distribution is a GP with a mean function μ_0 and kernel K_0 given by (1), and use \mathbb{E}_0 to indicate the expectation taken with respect to this distribution. We let $f_0^* = \min f(D)$ be the best point observed thus far.

We index quantities associated with the first evaluation remaining by 1, and the second by 2. We let X_1 indicate the set of points to be evaluated in the first stage. We let $f(X_1) = (f(x) : x \in X_1)$ indicate the corresponding vector of observed values and let $\min f(X_1)$ be the smallest value in this vector. We let x_2 indicate the single point observed in the second stage.

For each $i = 1, 2$, we define f_i to be smallest value observed by the end of stage i , so $f_1 = \min(f_0^*, f(X_1))$ and $f_2 = \min(f_1^*, f(x_2))$. We let μ_i be the mean function and K_i the kernel for the posterior mean given D and observations available at the end of stage i . We let \mathbb{E}_i indicate the expectation taken with respect to the corresponding Gaussian process.

The overall loss whose expected value we seek to minimize is f_2^* .

To find the optimal sampling strategy, we follow the dynamic programming principle. We write down the expected loss achievable at the start of the second stage, but supposing that evaluations at points in X_1 have already completed and the resulting values $f(X_1)$ observed. If we choose the final evaluation optimally, then the expected loss is $L_1 = \min_{x_2} \mathbb{E}_1[f_2^*]$. This posterior and thus also L_1 depends on X_1 and $f(X_1)$.

Following the derivation of the expected improvement Jones et al. [1998], we rewrite this as

$$L_1 = \min_{x_2} \mathbb{E}_1[f_1^* - (f_1^* - f(x_2))^+] = f_1^* - \max_{x_2} \mathbb{E}_1[(f_1^* - f(x_2))^+] = f_1^* - \max_{x_2} \text{EI}_1(x_2),$$

where $y^+ = \max(y, 0)$ is the positive part function and $\text{EI}_1(x)$ is the expected improvement criterion under the GP after the first evaluation has been performed:

$$\text{EI}_1(x) = \text{EI}(f_1^* - \mu_1(x_2), K_1(x_2, x_2)) \quad (2)$$

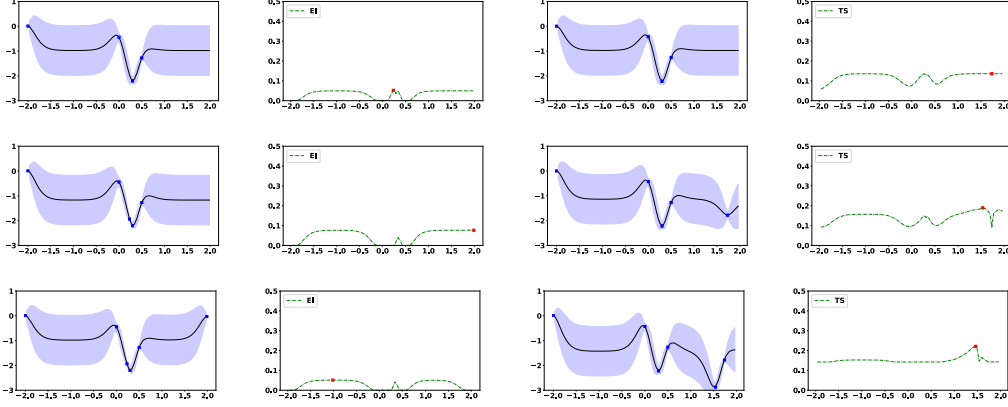


Figure 1: We demonstrate Two-Step (TS) and EI minimizing a 1-d synthetic function sampled from a GP. Each row shows the posterior on f (mean \pm one standard deviation) and the corresponding acquisition function, for EI (left) and TS (right). We plot progress over three iterations. On the first iteration, EI evaluates a point that refines an existing local optimum and could have provided a small one-step improvement, but provides little information of use in future evaluations. In contrast, TS explores more aggressively, which helps it identify a new global minimum in the next iteration.

where $\text{EI}(m, v) = m\Phi(m/\sqrt{v}) + \varphi(m/\sqrt{v})$ gives the expected improvement at a point where the difference between the best observed point and the mean is m , and the variance is v . Here, Φ is the standard normal cdf, and φ is the standard normal pdf.

With this expression for the value achievable at the start of the second stage, the expected value achieved at the start of the first stage is:

$$\begin{aligned} \mathbb{E}_0[L_1] &= \mathbb{E}_0 \left[f_1^* - \max_{x_2} \text{EI}_1(x_2) \right] = \mathbb{E}_0 \left[f_0^* - (f_0^* - \min f(X_1))^+ - \max_{x_2} \text{EI}_1(x_2) \right] \\ &= f_0^* - \text{EI}_0(X_1) - \mathbb{E}_0 \left[\max_{x_2} \text{EI}_1(x_2) \right], \end{aligned} \quad (3)$$

where $\text{EI}_0(X_1) = \mathbb{E}_0[(f_0^* - \min f(X_1))^+]$ is the multipoints expected improvement Ginsbourger et al. [2010] under the GP with mean μ_0 and kernel K_0 .

We define our two-step acquisition function to be

$$2\text{-OPT}(X_1) = \text{EI}_0(X_1) + \mathbb{E}_0 \left[\max_{x_2} \text{EI}_1(x_2) \right]. \quad (4)$$

Because f_0^* does not depend on where we evaluate next X_1 , finding the X_1 that minimizes (3) is equivalent to finding the value that maximizes Q .

Figure 1 illustrates the behavior of the two-step acquisition function, showing how it provides more exploration than EI.

2.3 Monte Carlo estimation of 2-OPT(\cdot)

2-OPT(X_1) cannot be computed in closed form. We can, however, estimate it using Monte Carlo. We first use the reparameterization trick [Wilson et al., 2018, L’Ecuyer, 1990] to write $f(X_1)$ as $\mu_0(X_1) + C_0(X_1)Z$, where Z is an q -dimensional independent standard normal random variable and $C_0(X_1)$ is the Cholesky decomposition of $K_0(X_1, X_1)$.

We assume that $K_0(X_1, X_1)$ is positive definite so $C_0(X_1)$ is of full rank.

Then, under (1), for generic x ,

$$\begin{aligned} \mu_1(x) &= \mu_0(x) + K_0(x, X_1)K_0(X_1, X_1)^{-1}(f(X_1) - \mu_0(X_1)) \\ &= \mu_0(x) + K_0(x, X_1)(C_0(X_1)C_0(X_1)^T)^{-1}C_0(X_1)Z \\ &= \mu_0(x) + \sigma_0(x, X_1)Z \end{aligned}$$

$$\begin{aligned}
K_1(x, x) &= K_0(x) - K_0(x, X_1)K_0(X_1, X_1)^{-1}K(X_1, x) \\
&= K_0(x) - K_0(x, X_1)(C_0(X_1)C_0(X_1)^T)^{-1}K(X_1, x) \\
&= K_0(x) - \sigma_0(x, X_1)\sigma_0(x, X_1)^T.
\end{aligned}$$

where $\sigma_0(x, X_1) = K_0(x, X_1)C_0(X_1)^{-1}$.

With this notation, we can write $\text{EI}_2(x_2)$ explicitly as

$$\text{EI}_2(x_2) = \text{EI}(\mu_0(x_2) + \sigma_0(x_2, X_1)Z, K_0(x_2) - \sigma_0(x_2, X_1)\sigma_0(x_2, X_1)^T) =: \Gamma(X_1, x_2, Z)$$

where we have defined a new function Γ for later convenience.

Thus, we can rewrite the 2-OPT acquisition function as $2\text{-OPT} = E_0[\widehat{2\text{-OPT}}(Z)]$ where

$$\begin{aligned}
\widehat{2\text{-OPT}}(Z) &= (f_0^* - \min f(X_1))^+ + \max_{x_2} \text{EI}_1(x_2) \\
&= \max(f_0^* - \mu_0(X_1) - C_0(X_1)Z)^+ + \max_{x_2} \Gamma(X_1, x_2, Z)
\end{aligned}$$

Then, to provide an unbiased estimate of $2\text{-OPT}(x_1)$, we sample Z and compute $\widehat{2\text{-OPT}}(Z)$ using a nonlinear global optimization routine to calculate the inner maximization involved. Averaging many such replications provides a strongly consistent estimator of $2\text{-OPT}(x_1)$.

Previous approaches [Osborne et al., 2009, Ginsbourger and Riche, 2010, González et al., 2016] use this or a similar simulation method to obtain an estimator of 2-OPT, and then use this estimator within a derivative-free optimization approach. This requires extensive computation because:

1. The nested optimization over x_2 is time-consuming, and must be done for each simulation;
2. Noise in the simulation requires either a noise-tolerant derivative-free optimization method that would typically require more iterations, or requires that the simulation be averaged often enough on each iteration to make noise negligible. This increases the number of simulations required to optimize accurately;
3. It does not leverage derivative information, causing optimization to require more iterations, especially as the dimension d of the search space or the batch size q grows.

3 Efficiently Optimizing 2-OPT

Here we describe a novel computational approach to optimizing 2-OPT which is substantially more efficient than previously proposed methods. Our approach includes two components: a novel simulation-based stochastic gradient estimator, which can be used within multistart stochastic gradient ascent; and variance reduction techniques that reduce the variance of this stochastic gradient estimator.

3.1 Unbiased Estimation of the Gradient of 2-OPT

We now show how to obtain an unbiased and strongly consistent estimator of the gradient of 2-OPT. The main idea is to exchange the expectation and the gradient operators

$$\begin{aligned}
\nabla 2\text{-OPT}(X_1) &= \mathbb{E}_0 \left[\nabla \widehat{2\text{-OPT}}(Z) \right] \\
&= \mathbb{E}_0 \left[\nabla \max(f_0^* - \mu_0(X_1) - C_0(X_1)Z)^+ + \nabla \max_{x_2} \Gamma(X_1, x_2, Z) \right] \\
&= \mathbb{E}_0 \left[\nabla \max(f_0^* - \mu_0(X_1) - C_0(X_1)Z)^+ + \nabla \Gamma(X_1, x_2^*, Z) \right]
\end{aligned}$$

where $x_2^* \in \arg \max_{x_2 \in \mathcal{A}} \Gamma(X_1, x_2, Z)$ is fixed and the last equation follows under some regularity conditions by the envelope theorem [Milgrom and Segal, 2002]. The following theorem shows this estimator of $\nabla 2\text{-OPT}$ is unbiased and strongly consistent. Its proof is in the supplement.

Theorem 1. We assume that the domain \mathcal{A} is compact, μ_0 is a constant and the kernel K_0 is continuously differentiable. For a given vector X_1 , we assume that $\max_{x_2 \in \mathcal{A}} \Gamma_n(X_1, x_2, Z)$ contains only one element almost surely. Let x_2^* be a global minimum in \mathcal{A} of $\Gamma(X_1, x_2, Z)$. We then have that

$$g(Z) := \nabla \max(f_0^* - \mu_0(X_1) - C_0(X_1)Z))^+ + \nabla \Gamma(X_1, x_2^*, Z) \quad (5)$$

exists almost surely and is an unbiased estimator of $\nabla 2\text{-OPT}(X_1)$.

We can use the previous theorem to obtain an unbiased estimator of the gradient of 2-OPT, which can be used within stochastic gradient ascent [Kushner and Yin, 2003] to maximize 2-OPT. Our following theorem shows that, under the right conditions, a stochastic gradient ascent algorithm converges almost surely to a critical point of 2-OPT. Its proof is in the supplement.

Theorem 2. Assume the conditions of Theorem 1 and that \mathcal{A} is a hyperrectangle. In addition, assume that we optimize 2-OPT using the stochastic gradient estimator ascent method with the stochastic gradient from Theorem 1. Furthermore, assume that the stepsize sequence $\{\epsilon_t : t = 0, 1, \dots\}$ of the SGA satisfies $\epsilon_t \rightarrow 0$, $\epsilon_t \geq 0$, $\sum_t \epsilon_t = \infty$ and $\sum_t \epsilon_t^2 < \infty$. Then the sequence of iterates produced by the stochastic gradient ascent algorithm converges almost surely to a connected set of stationary points of 2-OPT.

We then use stochastic gradient ascent with multiple restarts to find a collection of stationary points X_1 . We use Monte Carlo to evaluate 2-OPT for each of these stationary points and select as our maximizer of 2-OPT the point or batch of points with the largest estimated 2-OPT(X_1).

3.2 Variance reduction

We now describe variance reduction techniques that further improve computation time and accuracy.

Gauss-Hermite Quadrature (fully sequential setting) In the fully sequential setting where we propose one point at each iteration ($q = 1$), we use Gauss-Hermite quadrature [Liu and Pierce, 1994] to estimate 2-OPT(X_1) and its gradient. These quantities are both expectations over the 1-d standard Gaussian random variable Z . Gauss-Hermite quadrature estimates the expectation of a random variable $g(Z)$ can be approximated by a weighted sum $\sum_{i=1}^n w_i g(z_i)$ with well-chosen weights w_i and locations z_i . In practice, we find $n = 20$ accurately estimates 2-OPT(X_1) and its gradient.

Importance sampling (batch setting) In the batch setting, Gauss-Hermite quadrature scales poorly with batch size q since the number of weighted points required grows exponentially with the dimension over which we integrate, which is q . In the batch setting, we adopt another variance reduction technique: importance sampling [Asmussen and Glynn, 2007].

Recall that our Monte Carlo estimator of 2-OPT and its stochastic gradient estimator involve a sampled multipoints EI term $\max(f_0^* - \mu_0(X_1) - C_0(X_1)Z))^+$.

For high-dimensional test functions or after we have many function evaluations, most draws of Z result in this multipoints EI term taking a value of 0. This occurs when all components of $\mu_0(X_1) + C_0(X_1)Z$ are larger than f_0^* . For such Z , the derivative of this immediate improvement term is also 0. Also, for such Z , the second term in our Monte Carlo estimator of 2-OPT and its stochastic gradient, $\max_{x_2} \Gamma(X_1, x_2, Z)$ also tends to be small and have a small gradient.

As a result, when calculating the expected value of these samples of 2-OPT or its gradient, we include many 0s. This can make the variance of estimators based on averaging these estimators large relative to their expected value. This in turn makes gradient-based optimization and comparison using Monte Carlo estimates challenging.

To address this, we simulate Z from a multivariate Gaussian distribution with a larger standard deviation $v > 1$, calling it Z^v . This substantially increases the chance that the at least one component of $\mu_0(X_1) + C_0(X_1)Z$ will exceed f_0^* . In practice, we find $v = 3$ works well in test problems.

To compensate for sampling from a different distribution, we multiply by the likelihood ratio between the density for which we wish to calculate the expectation, which is the multivariate standard normal density, and the density from which Z^v was sampled. Letting $\varphi(\cdot; 0, vI)$ indicate the q -dimensional normal multivariate density with mean 0 and covariance matrix vI , this likelihood ratio is $\varphi(Z^v; 0, I)/\varphi(Z^v; 0, vI)$.

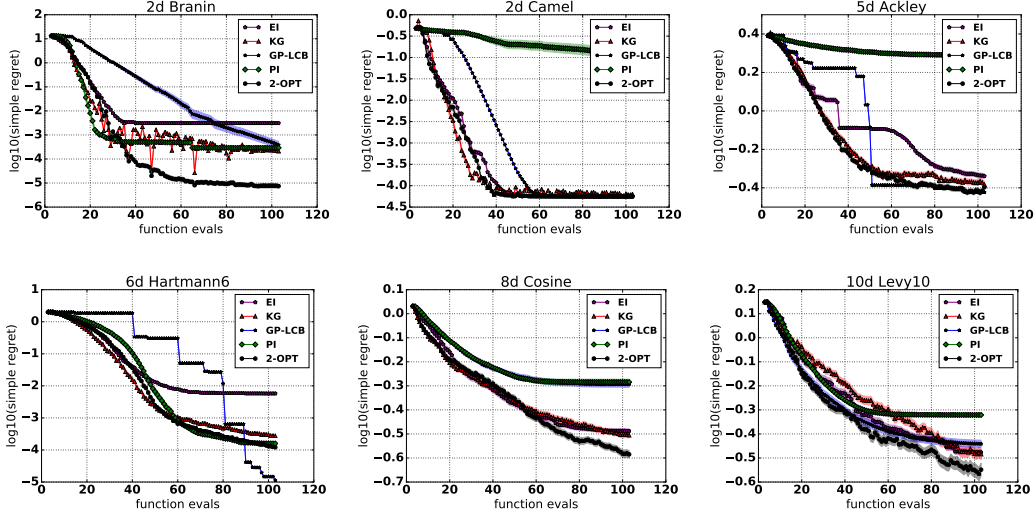


Figure 2: Benchmarks of 2-OPT with common one-step heuristics: EI, KG and GP-LCB on six common synthetic functions. 2-OPT outperforms the competitors on 5 out of 6 test functions, although some of the one-step algorithms are known to be highly effective on these functions. On the 10-d Levy10 test function, 2-OPT falls behind EI a little in the end but leads all the competitors in the most time of the progress.




Figure 3: HPOlib and ATO benchmarks: suite: the logistic regression benchmark (left) and the SVM benchmark (right). 2-OPT is competitive with (outperforms) the best of the competitors in both benchmarks. The ATO benchmark: 2-OPT outperforms EI slightly, both outperforms KG clearly. However, all the algorithms converge with almost the same performance.

230 The resulting unbiased estimators of 2-OPT and its gradients are, respectively,
 231 $\widehat{2\text{-OPT}}(Z^v)\varphi(Z^v; 0, I)/\varphi(Z^v; 0, vI)$ and $g(Z^v)\varphi(Z^v; 0, I)/\varphi(Z^v; 0, vI)$.

232 4 Numerical experiments

233 We test our algorithms on common synthetic functions and widely-benchmarked real-world. We
 234 compare with acquisition functions widely used in practice including GP-LCB [Srinivas et al., 2010],
 235 EI [Snoek et al., 2012] and KG [Wu and Frazier, 2016] and multi-step lookahead methods from Lam
 236 et al. [2016] and González et al. [2016]. 2-OPT outperforms these competitors on most benchmarks
 237 and is at least comparable across all benchmarks. The computational time of 2-OPT varies from
 238 seconds to several minutes in all our experiments. In comparison, the algorithm in [Lam, 2018] takes 10 minutes
 239 to 1 hour for each suggestion [Lam, 2018]. 2-OPT’s strong empirical performance together with a
 240 supporting fast computational method unlocks the value of two-step lookahead in practice.

241 Before describing results, we highlight experimental details. Following Snoek et al. [2012], we use
 242 a constant mean prior and the ARD Matérn 5/2 kernel. We integrate over GP hyperparameters by
 243 sampling 16 sets of values using the emcee package [Foreman-Mackey et al., 2013]. We will make
 244 our implementation and the code to replicate experiments public upon publication.

245 **Synthetic functions, compared with one-step methods.** First, we test our algorithm along with the
 246 benchmarks on 6 well-known synthetic test functions chosen from Bingham [2015] ranging from
 247 2d to 10d: 2d Branin, 2d Camel, 5d Ackley5, 6d Hartmann6, 8d Cosine and 10d Levy function. We

| Function name | | PI | EI | UCB | PES | GLASSES | R-4-9 | R-4-10 | R-5-9 | R-5-10 | 2-OPT |
|-----------------|--------|------|-------------|------|-------------|------------------|-------------|--------|-------|--------|--------------|
| Branin-Hoo | Mean | .847 | .818 | .848 | .861 | .846 | .904 | .898 | .887 | .903 | .9995 |
| | Median | .922 | .909 | .910 | .983 | .909 | .959 | .943 | .921 | .950 | .9994 |
| Goldstein-Price | Mean | .873 | .866 | .733 | .819 | .782 | .895 | .784 | .861 | .743 | .9651 |
| | Median | .983 | .981 | .899 | .987 | .919 | .991 | .985 | .989 | .928 | .9911 |
| Griewank | Mean | .827 | .884 | .913 | .972 | 1.0 ¹ | .882 | .885 | .930 | .867 | .9321 |
| | Median | .904 | .953 | .970 | .987 | 1.0 ¹ | .967 | .962 | .960 | .954 | .9801 |
| 6-hump Camel | Mean | .850 | .887 | .817 | .664 | .776 | .860 | .825 | .793 | .803 | .9010 |
| | Median | .893 | .970 | .915 | .801 | .941 | .926 | .900 | .941 | .907 | .9651 |

Table 1: Performance of our two-step acquisition function (2-OPT) on test functions compared with non-myopic and other benchmark algorithms originally reported in Lam et al. [2016]. Each value reported is the “gap”: the ratio of the overall improvement obtained by the algorithm to the improvement possible by a globally optimal solution. A gap of 1 represents finding the optimal solution; 0 represents no improvement in solution quality. The best gap appears in boldface.

initiate our algorithms by randomly sampling 3 points from a Latin hypercube design and then start the Bayesian optimization iterative process. Figure 2 reports the mean and the standard deviation of the base 10 logarithm of the immediate regret by running 100 random initializations.

Synthetic functions, compared with multi-step methods. To compare with non-myopic algorithms proposed in González et al. [2016] and Lam et al. [2016], we replicate the experimental settings in Lam et al. [2016] and add the performances of our algorithm to their Table 2. We report the results in Table 1. GLASSES was proposed in González et al. [2016] and the four columns R-4-9, R-4-10, R-5-9, and R-5-10 are algorithm variants proposed in Lam et al. [2016].

Values reported are the “gap” [Huang et al., 2006], which is the ratio of the improvement obtained by the solution reported by the algorithm \hat{x}_N given its full sampling budget N to the improvement possible by a globally optimal solution. The improvement is measured relative to the best solution found in the initial stage of samples performed uniformly at random. Letting \hat{x}_N be the best solution found by the algorithm and \hat{x}_0 be the best solution found in the initial stage, the gap is $G = (f(\hat{x}_0) - f(\hat{x}_N)) / (f(\hat{x}_0) - \min_x f(x))$. A gap of 1 indicates that the algorithm found a globally optimal solution, while 0 indicates no improvement.

2-OPT is best in 5 out of 8 problems (tied for best on one of these problems), and second-best in the remaining 3. It outperforms or ties the non-myopic competitors on all problem instances.

HPOlib hyperparameter tuning benchmarks Snoek et al. [2012] provides hyperparameter tuning benchmarks for BayesOpt, further summarized in Eggenberger et al. [2013] and packaged into the HPOlib library. We benchmark on the two most widely used test problems there: logistic regression and SVM. On both problems, 2-OPT performs comparably to the best of the competitors, with 2-OPT and EI slightly outperforming KG on logistic regression.

Assemble to order (ATO) benchmark The assemble-to-order (ATO) benchmark [Hong and Nelson, 2006, Poloczek et al., 2017] is a reinforcement learning problem where the goal is to optimize an 8-dimensional inventory target vector to maximize profit in a business setting. 2-OPT provides a substantial benefit over competitors from the start and remains best over the whole process. After 40 iterations, EI catches 2-OPT, while KG lags both EI and 2-OPT until iteration 100 where all the algorithms converge with comparable performance.

5 Conclusions

In this article, we propose the first computationally efficient two-step lookahead BayesOpt algorithm. The algorithm comes in both sequential and batch forms, and reduces the computational time compared to previous proposals with increased performance. In experiments, we find that two-step lookahead provides additional value compared to several one-step lookahead heuristics.

¹As reported in Lam et al. [2016], GLASSES achieves gap=1 on Griewank because an optimizer it uses arbitrarily evaluates at the origin, which happens to be a global minimizer. Following Lam et al. [2016], we exclude these results.

References

- S. Asmussen and P. W. Glynn. *Stochastic simulation: algorithms and analysis*, volume 57. Springer Science & Business Media, 2007.
- D. Bingham. Optimization test problems. <http://www.sfu.ca/~ssurjano/optimization.html>, 2015.
- E. Brochu, V. M. Cora, and N. De Freitas. A tutorial on bayesian optimization of expensive cost functions, with application to active user modeling and hierarchical reinforcement learning. *arXiv preprint arXiv:1012.2599*, 2010.
- K. Eggenberger, M. Feurer, F. Hutter, J. Bergstra, J. Snoek, H. Hoos, and K. Leyton-Brown. Towards an empirical foundation for assessing bayesian optimization of hyperparameters. In *NIPS workshop on Bayesian Optimization in Theory and Practice*, volume 10, page 3, 2013.
- D. Foreman-Mackey, D. W. Hogg, D. Lang, and J. Goodman. emcee: the mcmc hammer. *Publications of the Astronomical Society of the Pacific*, 125(925):306, 2013.
- A. Forrester, A. Sobester, and A. Keane. *Engineering design via surrogate modelling: a practical guide*. John Wiley & Sons, 2008.
- D. Ginsbourger and R. Riche. Towards gaussian process-based optimization with finite time horizon. *mODa 9—Advances in Model-Oriented Design and Analysis*, pages 89–96, 2010.
- D. Ginsbourger, R. Le Riche, and L. Carraro. Kriging is well-suited to parallelize optimization. In *Computational Intelligence in Expensive Optimization Problems*, pages 131–162. Springer, 2010.
- J. González, M. Osborne, and N. Lawrence. Glasses: Relieving the myopia of bayesian optimisation. In *Artificial Intelligence and Statistics*, pages 790–799, 2016.
- P. Heidelberger, X.-R. Cao, M. A. Zazanis, and R. Suri. Convergence properties of infinitesimal perturbation analysis estimates. *Management Science*, 34(11):1281–1302, 1988.
- J. M. Hernández-Lobato, M. W. Hoffman, and Z. Ghahramani. Predictive entropy search for efficient global optimization of black-box functions. In *Advances in Neural Information Processing Systems*, pages 918–926, 2014.
- L. J. Hong and B. L. Nelson. Discrete optimization via simulation using compass. *Operations Research*, 54(1):115–129, 2006.
- D. Huang, T. T. Allen, W. I. Notz, and N. Zeng. Global optimization of stochastic black-box systems via sequential kriging meta-models. *Journal of global optimization*, 34(3):441–466, 2006.
- D. R. Jones, M. Schonlau, and W. J. Welch. Efficient global optimization of expensive black-box functions. *Journal of Global optimization*, 13(4):455–492, 1998.
- H. Kushner and G. G. Yin. *Stochastic approximation and recursive algorithms and applications*, volume 35. Springer Science & Business Media, 2003.
- H. J. Kushner. A new method of locating the maximum point of an arbitrary multipeak curve in the presence of noise. *Journal of Fluids Engineering*, 86(1):97–106, 1964.
- R. Lam. Personal communication, 2018.
- R. Lam, K. Willcox, and D. H. Wolpert. Bayesian optimization with a finite budget: An approximate dynamic programming approach. In *Advances in Neural Information Processing Systems*, pages 883–891, 2016.
- P. L’Ecuyer. A unified view of the IPA, SF, and LR gradient estimation techniques. *Management Science*, 36(11):1364–1383, 1990.
- Q. Liu and D. A. Pierce. A note on gauss—hermite quadrature. *Biometrika*, 81(3):624–629, 1994.
- P. Milgrom and I. Segal. Envelope theorems for arbitrary choice sets. *Econometrica*, 70(2):583–601, 2002.

326 M. A. Osborne, R. Garnett, and S. J. Roberts. Gaussian processes for global optimization. In *3rd*
327 *international conference on learning and intelligent optimization (LION3)*, pages 1–15. Citeseer,
328 2009.

329 M. Poloczek, J. Wang, and P. I. Frazier. Multi-information source optimization. In *Advances in Neural*
330 *Information Processing Systems*, 2017. Accepted for publication. ArXiv preprint 1603.00389.

331 W. B. Powell. *Approximate Dynamic Programming: Solving the curses of dimensionality*, volume
332 703. John Wiley & Sons, 2007.

333 C. E. Rasmussen and C. K. I. Williams. *Gaussian Processes for Machine Learning*. MIT Press, 2006.
334 ISBN ISBN 0-262-18253-X.

335 J. Snoek, H. Larochelle, and R. P. Adams. Practical bayesian optimization of machine learning
336 algorithms. In *Advances in neural information processing systems*, pages 2951–2959, 2012.

337 N. Srinivas, A. Krause, M. Seeger, and S. M. Kakade. Gaussian process optimization in the bandit
338 setting: No regret and experimental design. In *ICML*, pages 1015–1022, 2010.

339 B. S. Thomson, J. B. Bruckner, and A. M. Bruckner. *Elementary real analysis*. ClassicalRealAnalysis.
340 com, 2008.

341 Z. Wang and S. Jegelka. Max-value entropy search for efficient bayesian optimization. *arXiv preprint*
342 *arXiv:1703.01968*, 2017.

343 J. Wilson, F. Hutter, and M. Deisenroth. Maximizing acquisition functions for bayesian optimization.
344 In *Advances in Neural Information Processing Systems*, pages 9884–9895, 2018.

345 J. Wu and P. Frazier. The parallel knowledge gradient method for batch bayesian optimization. In
346 *Advances in Neural Information Processing Systems*, pages 3126–3134, 2016.

347 J. Wu, M. Poloczek, A. G. Wilson, and P. I. Frazier. Bayesian optimization with gradients. In
348 *Advances in Neural Information Processing Systems*, 2017. Accepted for publication. ArXiv
349 preprint 1703.04389.